

1 February 4th, 2021

1.1 Efficient Gaussian Elimination

Last time, we introduced a memory efficient way to perform Gaussian Elimination by reusing the inputs. However, this implementation is **not** the most time efficient. To do this, we need to consider the computer architecture. Recall that computers have a few different layers of memory:

- CPU register
- Cache
- RAM
- Disk

The lower the level, the larger the memory but communicating is slower. As such, to improve the efficiency, we should minimize the communication between the different levels of memory.

Remark 1.1 — This principle of memory management should always be considered when performing matrix multiplication.

Note that we can reformulate the previous algorithms into matrix operations. First we perform a vector scaling, then we perform a rank 1 update, and then a matrix vector product. Since these are very basic operations, there are many very efficient implementations of these **Basic Linear Algebra Subroutines (BLAS)**. There are a few different levels of these BLAS.

BLAS1 : Vector operations, e.g.

$$\beta = x^T y, y = \beta x + y, \quad \text{where } x, y \in \mathbb{R}^n, \beta \in \mathbb{R}$$

BLAS2 : Matrix-Vector operations

$$y = Ax + y, \quad \text{where } A \in \mathbb{R}^{m \times n}, x, y \text{ are vectors}$$

BLAS3 : Matrix-Matrix operations

$$C = AB + C, \quad \text{where } A, B, C \text{ are matrices}$$

These BLAS are implemented to be optimized on the given computer architecture.

Remark 1.2 — Since BLAS3 can be constructed using BLAS2, etc. in terms of the efficiency, we have:

$$\text{BLAS3} > \text{BLAS2} > \text{BLAS1}$$

where efficiency is the time needed

Remark 1.3 — BLAS will come with the CPU.

AMD Core Math Library (ACML)

Intel Math Kernel Lib (MKL)

Replacing the algorithm with BLAS, we have:

Algorithm 1 Gaussian Elimination with BLAS

```

1: for  $k = 1 : n - 1$  do
2:    $A(k+1:n, k) = A(k+1:n, k)/A(k, k)$ 
3:    $A(k+1:n, k+1:n) = A(k+1:n, k+1:n) - A(k+1:n, k)A(k, k+1:n)$ 
4:    $b(k+1:n, 1) = b(k+1:n, 1) - b(k, 1)A(k+1:n, k)$ 
5: end for
6: for  $k = n : -1 : 1$  do
7:    $b(k) = b(k) - A(k, k+1:n)b(k+1:n)$ 
8:    $b(k) = b(k)/A(k, k)$ 
9: end for

```

1.2 Complexity of Gaussian Elimination

For this, we want to see how many scalar operations are needed. Counting the cost, we have:

$$\begin{aligned}
& \sum_{k=1}^{n-1} [(n-k) + 2(n-k)^2 + 2(n-k)] + \sum_{k=1}^n [2(n-k) + 1] \\
& \sum_{k=1}^{n-1} [(n-k) + 2(n-k)^2 + 2(n-k)] + \sum_{k=1}^n [2(n-k) + 1] \\
& = 3 \sum_{k=1}^{n-1} (n-k) + 2 \sum_{k=1}^{n-1} (n-k)^2 + 2 \sum_{k=1}^{n-1} k + \sum_{k=1}^n 1 \\
& = \frac{2}{3}n^3 + \frac{3}{2}n^2 - \frac{1}{3}n \\
& = \frac{2}{3}n^3 + O(n^2) \\
& = O(n^3).
\end{aligned}$$

Theorem 1.4

The computational complexity of Gaussian Elimination is $O(n^3)$.

1.3 LU Decomposition view of Gaussian Elimination

Remember we can express Gaussian Elimination as:

$$L_{n-1} \dots L_2 L_1 A x = L_{n-1} \dots L_1 b$$

Note that $L_{n-1} \dots L_2 L_1 A = A^{(n)}$ which is upper triangular. Let us denote this by:

$$U = \begin{bmatrix} u_{11} & u_{12} & \dots & u_{1n} \\ & u_{22} & \dots & u_{2n} \\ & & \ddots & \vdots \\ & & & u_{nn} \end{bmatrix} \in \mathbb{R}^{n \times n}$$

Note that since each L_i is a Gaussian transformation which is invertible, we have:

$$A = L_1^{-1} L_2^{-1} \dots L_{n-1}^{-1} U$$

Theorem 1.5

$$L_k^{-1} = I + \ell_k e_k^T$$

Proof.

$$\begin{aligned} & (I - \ell_k e_k^T)(I + \ell_k e_k^T) \\ &= I - \ell_k e_k^T + \ell_k e_k^T - \ell_k e_k^T \ell_k e_k^T \\ &= I. \end{aligned}$$

□

Theorem 1.6

$$L_1^{-1} L_2^{-1} \dots L_{n-1}^{-1} = I + \sum_{i=1}^{n-1} \ell_i e_i^T$$

which is lower triangular.

Proof. Direct computation. □

Theorem 1.7

Gaussian Elimination gives:

$$A = LU$$

where L is unit lower triangular, and U is upper triangular. This is called the **LU Decomposition**.

Remark 1.8 — Note that in the memory efficient implementation of GE, we are computing L and U and it is stored in the original matrix.

With this, we have:

$$Ax = b \iff LUx = b \iff \begin{cases} Ly = b \\ Ux = y \end{cases}$$

Theorem 1.9

The computation cost is:
$$\begin{cases} \text{LU decomposition} & \frac{2}{3}n^3 + O(n^2) \\ \text{Forward Substitution} & O(n^2) \\ \text{Backward Substitution} & O(n^2) \end{cases}$$

Remark 1.10 — If we use the same coefficient matrix, we only need to compute the decomposition once.

This LU decomposition view of Gaussian Elimination has a number of advantages:

1. If we want to solve $Ax_i = b_i$ for $i = 1, \dots, m$, then we only need to do LU decomposition once
2. LU decomposition can be used for other matrix computation tasks. e.g.
 - (a) $\det(A) = \det(LU) = \det(L)\det(U) = u_{11}u_{22} \dots u_{nn}$
 - (b) $A^{-1} = U^{-1}L^{-1}$